

OS 161:

Address Space & Memory Management



Politecnico
di Torino

Department of Control and
Computer Engineering



System Programming - Sarah Azimi

CAD & Reliability Group
DAUIN- Politecnico di Torino

Today:

- Memory Management in OS161
 - Kmalloc for kernel side
 - Dumbvm for user side
- MIPS Virtual Address Space
- Kernel Loader
- Allocating Memory

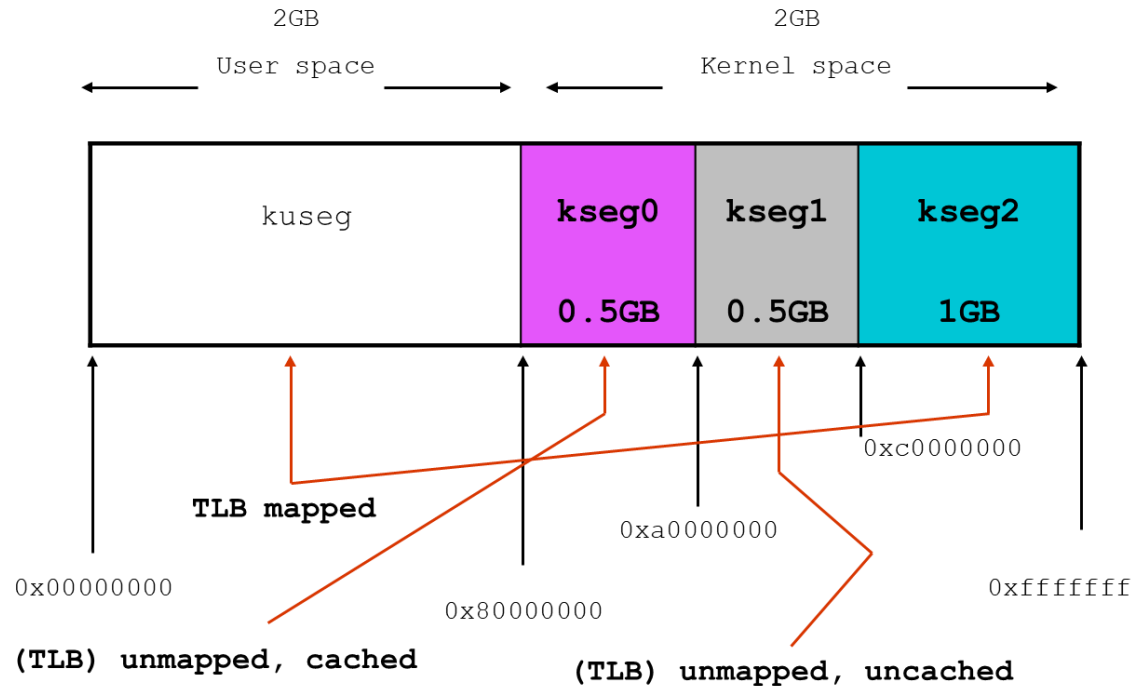
Dumbvm and kmalloc

- Memory management of OS161 consists of two elements:
 - Kmalloc allows memory allocation of kernel side
 - Dumbvm memory manager user side
- Memory manager of OS161 use **Contiguous allocation**
- However, it also uses pagination (minimum amount of available memory is page(s))

Dumbvm and kmalloc

- Allocation of memory is done in two levels:
 - `getppages` (`dumbvm.c`): calls `ram_stealmem` (in mutual exclusion) (more generic abstraction)
 - `ram_stealmem` (`ram.c`): allocates contiguous RAM starting at `firstpaddr`, that is increased (containing codes architecture dependent)
- Allocator is common to (for both user and kernel)
 - User memory: **`as_prepare_load`** calls **`getppages`** for 2 user segments and a stack
 - Dynamic kernel memory: **`kmalloc`** is based on **`alloc_kpages`**, that calls **`getppages`**

MIPS VIRTUAL ADDRESS SPACE



Mips maps the kernel of OS in the logical memory space of the process. So the process sees the user and also kernel in its memory space. (When boot of os161 is done, the kernel is mapped in a physical address that then later, through tlb, it is mapped in this logical memory of the process)

The kuseg is for user memory.

Kseg0 is for kernel, kseg01 I/O devices, and kseg02 is not used.

Kernel loader (sys161: start.S)

Logical addr. (KSEG0)

0x80000000

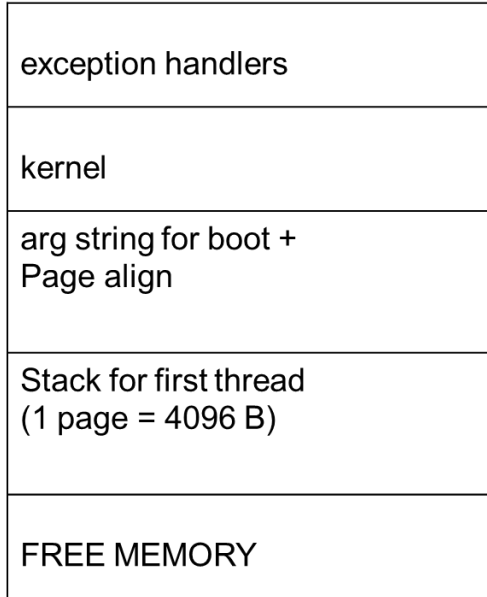
0x80000200

0x80039d54 (`_end`)

0x8003a000 (P)

0x8003b000 (P+1000)

0x80100000



ramsize (es. 1MB: sys161.conf)

Physical addr.

0x0

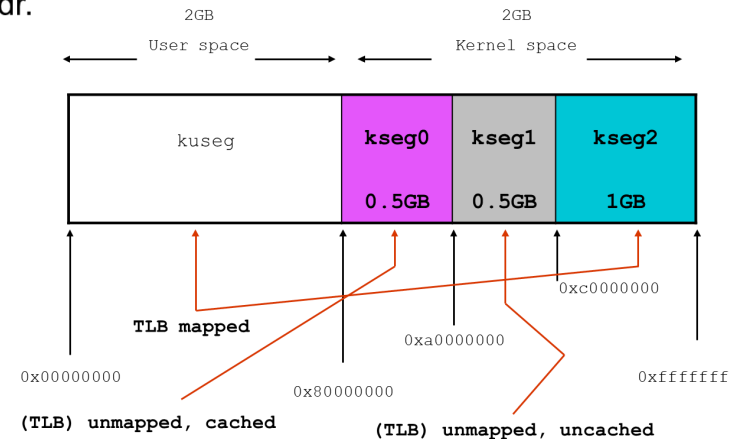
0x200

0x39d54

0x3a000

0x3b000

0x100000



What happens during the boot time of Os161?

Using part of the memory for loading the kernel of OS161, starting from address 0. Starting from exception handlers, kernel, the memory space for the argument of command line of kernel (it is empty right now), stack for the first process that will be created, and free physical memory).

Kernel loader (sys161: start.S)

Logical addr. (KSEG0)		Physical addr.
0x80000000		0x0
	exception handlers	
0x80000200		0x200
	kernel	
0x80039d54 (_end)		0x39d54
	arg string for boot + Page align	
0x8003a000 (P)		0x3a000
	Stack for first thread (1 page = 4096 B)	
0x8003b000 (firstfree)		0x3b000 (firstpaddr)
	FREE MEMORY	
0x80100000	ramsize (es. 1MB: sys161.conf)	0x100000

At the end of the boot phase, OS saves the first available physical address (Free Memory)

Dumbvm

Logical addr. (KSEG0)

0x80000000

0x8003b000

(firstfree)

0x80100000

kernel

DUMBVM

ram_stealmem

ramsize

Physical addr.

0x0

0x3b000

(firstpaddr)

0x100000

Complete layout of physical memory:

Starting from the first free address (first physical address), arriving to the dimension of the memory.

The first free address depends on the size of the kernel while the last depends on the ram size.

ram_bootstrap

```
void
ram_bootstrap(void) {
    /* Get size of RAM. */
    size_t ramsize = mainbus_ramsize();
    if (ramsize > 512*1024*1024) {
        ramsize = 512*1024*1024;
    }
    lastpaddr = ramsize;
    /* Get first free virtual address from where
       start.S saved it. Convert to physical address. */
    firstpaddr = firstfree - MIPS_KSEG0;
}
```

During the boot, the **ram_bootstrap** is called which is initializing different parameters such as **ramsize** using function **mainbus_ramsize**, saving **lastpaddr** as the ram size, and calculating the first free address with respect to KSEG0.

ram_stealmem (kern/arch/mips/vm/ram.c)

```
paddr_t ram_stealmem(unsigned long npages) {
    paddr_t paddr;
    size_t size = npages * PAGE_SIZE;
    if(firstpaddr + size > lastpaddr) {
        return 0;
    }
    paddr = firstpaddr;
    firstpaddr += size;
    return paddr;
}
```

In **ram.c**, the most important function is **ram_stealmem** that allows to ask for memory.

Npages is the number of pages to ask the OS.

Calculating the size as the multiply of number of pages and page size.

ram_stealmem (kern/arch/mips/vm/ram.c)

```
paddr_t ram_stealmem(unsigned long npages) {  
    paddr_t paddr;  
    size_t size = npages * PAGE_SIZE;  
    if(firstpaddr + size > lastpaddr) {  
        return 0;  
    }  
    paddr = firstpaddr;  
    firstpaddr += size;  
    return paddr;  
}
```

Checking whether the first free address + the requested size is more than the last physical address. If yes, returning 0 which shows failing in memory allocation

ram_stealmem (kern/arch/mips/vm/ram.c)

```
paddr_t ram_stealmem(unsigned long npages) {  
    paddr_t paddr;  
    size_t size = npages * PAGE_SIZE;  
    if(firstpaddr + size > lastpaddr) {  
        return 0;  
    }  
    paddr = firstpaddr;  
    firstpaddr += size;  
    return paddr;  
}
```

Checking whether the first free address + the requested size is more than the last physical address. If not, successful memory allocation, returning the address of the first available physical free memory.

Ram.c is regarding the architecture dependent part. For example, for changing from continues allocation to another one, this module should be changed.

Getppages (kern/arch/mips/vm/ram.c)

```
static paddr_t  
getppages(unsigned long npages) {  
    paddr_t addr;  
  
    spinlock_acquire(&stealmem_lock);  
    addr = ram_stealmem(npages);  
    spinlock_release(&stealmem_lock);  
    return addr;  
}
```

Getppages is for asking numbers of physical pages, asking for page number and returning pointer to the zone of the memory.

Getppages (kern/arch/mips/vm/ram.c)

```
static paddr_t  
getppages(unsigned long npages) {  
    paddr_t addr;  
  
    spinlock_acquire(&stealmem_lock);  
    addr = ram_stealmem(npages);  
    spinlock_release(&stealmem_lock);  
    return addr;  
}
```

Spinlock allows the mutual exclusion in case of multiple process, to avoid corrupting of ram_stealmem function.

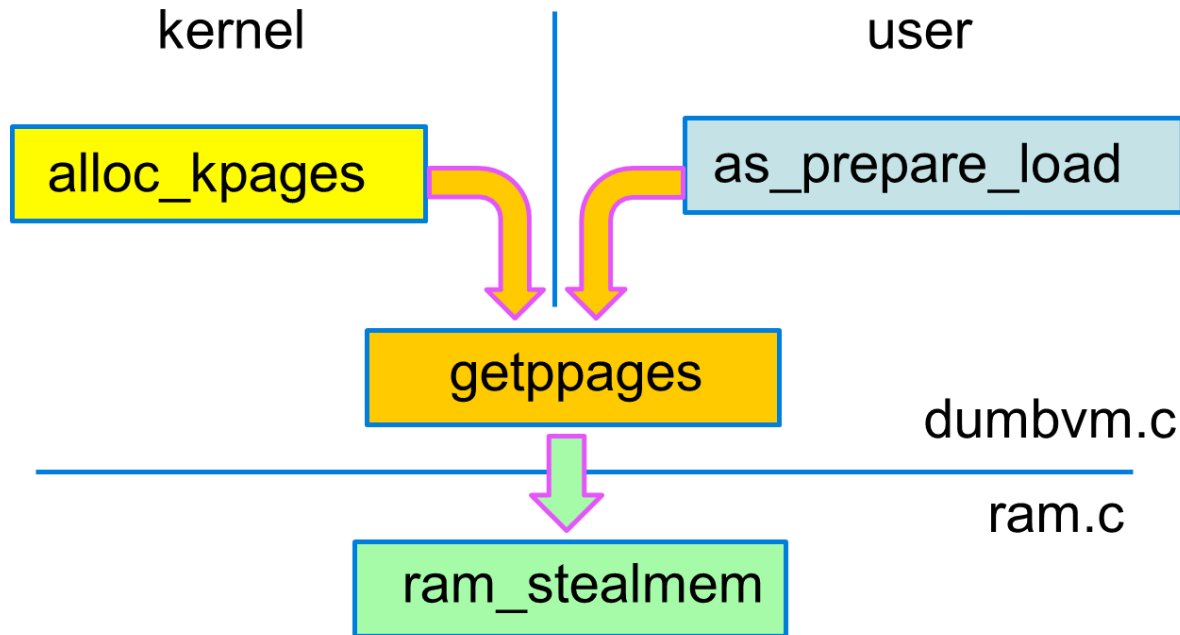
Getppages (kern/arch/mips/vm/ram.c)

```
static paddr_t  
getppages(unsigned long npages) {  
    paddr_t addr;  
  
    spinlock_acquire(&stealmem_lock);  
  
    addr = ram_stealmem(npages);  
  
    spinlock_release(&stealmem_lock);  
  
    return addr;  
}
```



Internal (dumbvm) function

Dumbvm.c (alloc)




Memory manager architecture of OS161:

ram_stealmem for memory management at low level (architecture dependent)

getppages for memory management in dumbvm.c

Allocation function for user side (**as_prepare_load**) and kernel side (**alloc_kpages**)

Dumbvm.c – Not implemented >> to do



The function for allocation of memory is implemented BUT the functions for releasing memory is not implemented.

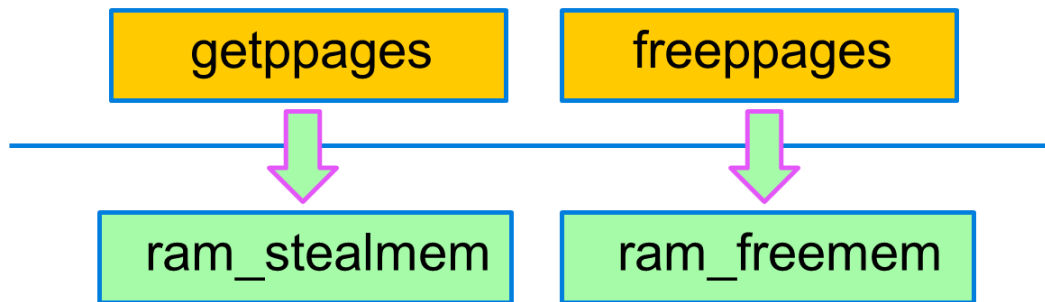
free_kpages and **as_destroy** (equivalent of **alloc_kpages** and **as_prepare_load**) is defined but EMPTY.

Your Job to solve this issue.

De-alloc (free) in ram.c

(solution A1)

- `freepages` just an interface to `ram_freemem`
- Data structure (free-list or bitmap) and **memory management in ram.c**

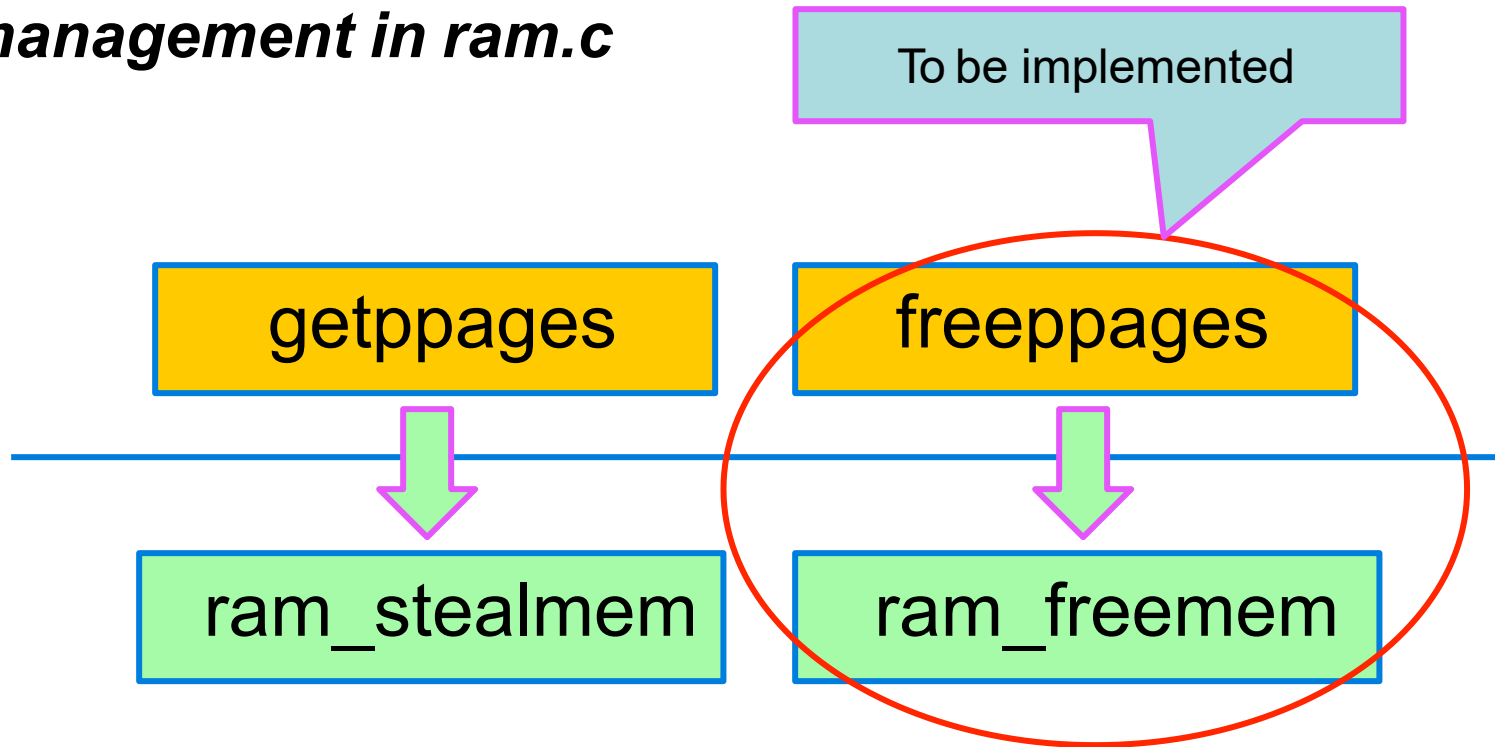


Defining the equivalent system of **ram_stealmem** and **getppages**,
Defining a **freepages** that we be called by **free_kpages** and **as_destroy** (high level)
that will call **ram_freemem** for freeing memory.

De-alloc (free) in ram.c

(solution A1)


- freeppages just an interface to ram_freemem
- Data structure (free-list or bitmap) and **memory management in ram.c**



De-alloc (free) in ram.c

(solution A2)

- Memory not returned to RAM
- Data structure (free-list o bitmap) and **memory management in dumbvm.c**
- Freeppages coordinates with getppages




Managing the liberation of memory at low level but just at the dumbvm.c level. Everytime that we are asking for memory allocation, before asking ram_stealmem, looking if the memory page that has been already allocated, is free now.

Vantages: working with one system call

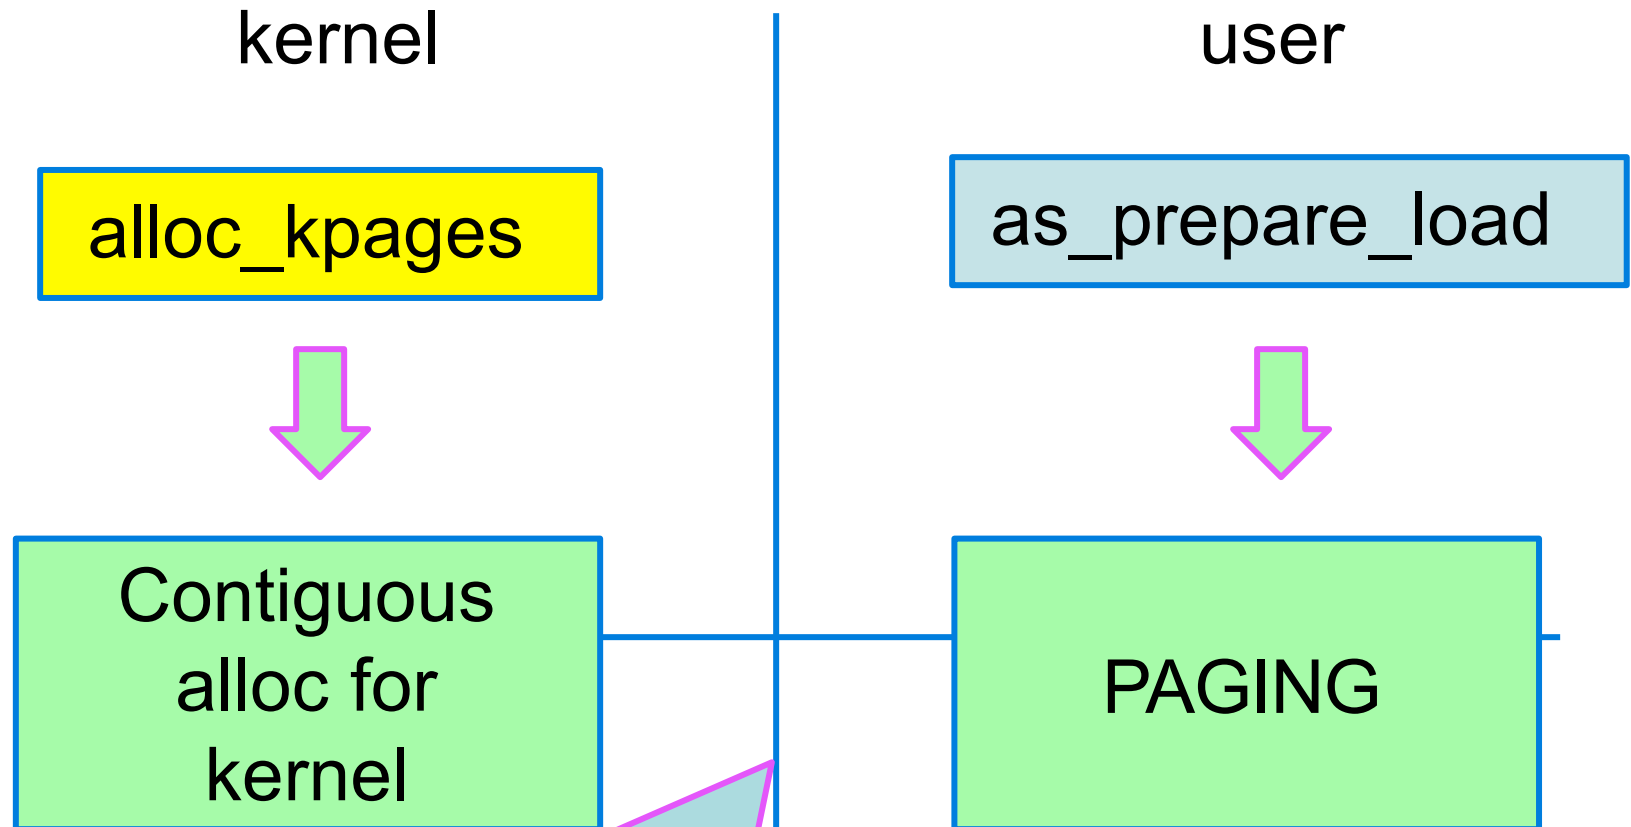
Paging in user space

(solution B)



Paging in user space

(solution B)



Two allocators:

- Pagine for user processes
- Contiguous memory for kernel

Proposed Solution

de-alloc in dumbvm (sol. A2)

- Contiguous allocation(by pages) common to kernel and user
- Allocator in dumbvm: keep track (using a bitmap) of previously freed pages. In order to alloc
 - When calling **getppages**, first search among (previously) freed pages (**an interval of contiguous free pages**)
 - **If not found, call ram_stealmem**
- Bitmap implemented as an array of char (for simplicity)
 - `freeRamFrames[i] = 1/0` (free/alloated): free=**FREED!** (by **freeppages**)
- In order to free we need to know
 - Pointer (or index) to first page in interval
 - Size, i.e. number of (contiguous) pages to free
- We need a table to store sizes (number of pages in allocated intervals) for each alloc performatd
 - `void free_kpages(vaddr_t addr) :` table needed as only pointer passed
 - `void as_destroy(struct addrspace *as) :` table not needed ad size is stored in address space
- `allocSize[i] = /* numbeer of pages allocated starting at i-th */`

Global Variables and Test Function

```
static struct spinlock freemem_lock = SPINLOCK_INITIALIZER;

static unsigned char *freeRamFrames = NULL;
static unsigned long *allocSize = NULL;
static int nRamFrames= 0;

static int allocTableActive = 0;

static int isTableActive () {
    int active;
    spinlock_acquire(&freemem_lock);
    active = allocTableActive;
    spinlock_release(&freemem_lock);
    return active;
}
```


Global Variables and Test Function

```
static struct spinlock freemem_lock = SPINLOCK_INITIALIZER;

static unsigned char *freeRamFrames = NULL;
static unsigned long *allocSize = NULL;
static int nRamFrames= 0;

static int allocTableActive = 0;



static int
    int active;
    spinlock_t lock;
    active = allocTableActive;
    spinlock_release(&freemem_lock);
    return active;
}
```

Dynamic arrays as RAM size known
at Boot (depends on da sys161.conf)
Alternative: over-dimensioned static arrays!

Dumbvm

Never freed

freeRamFrames



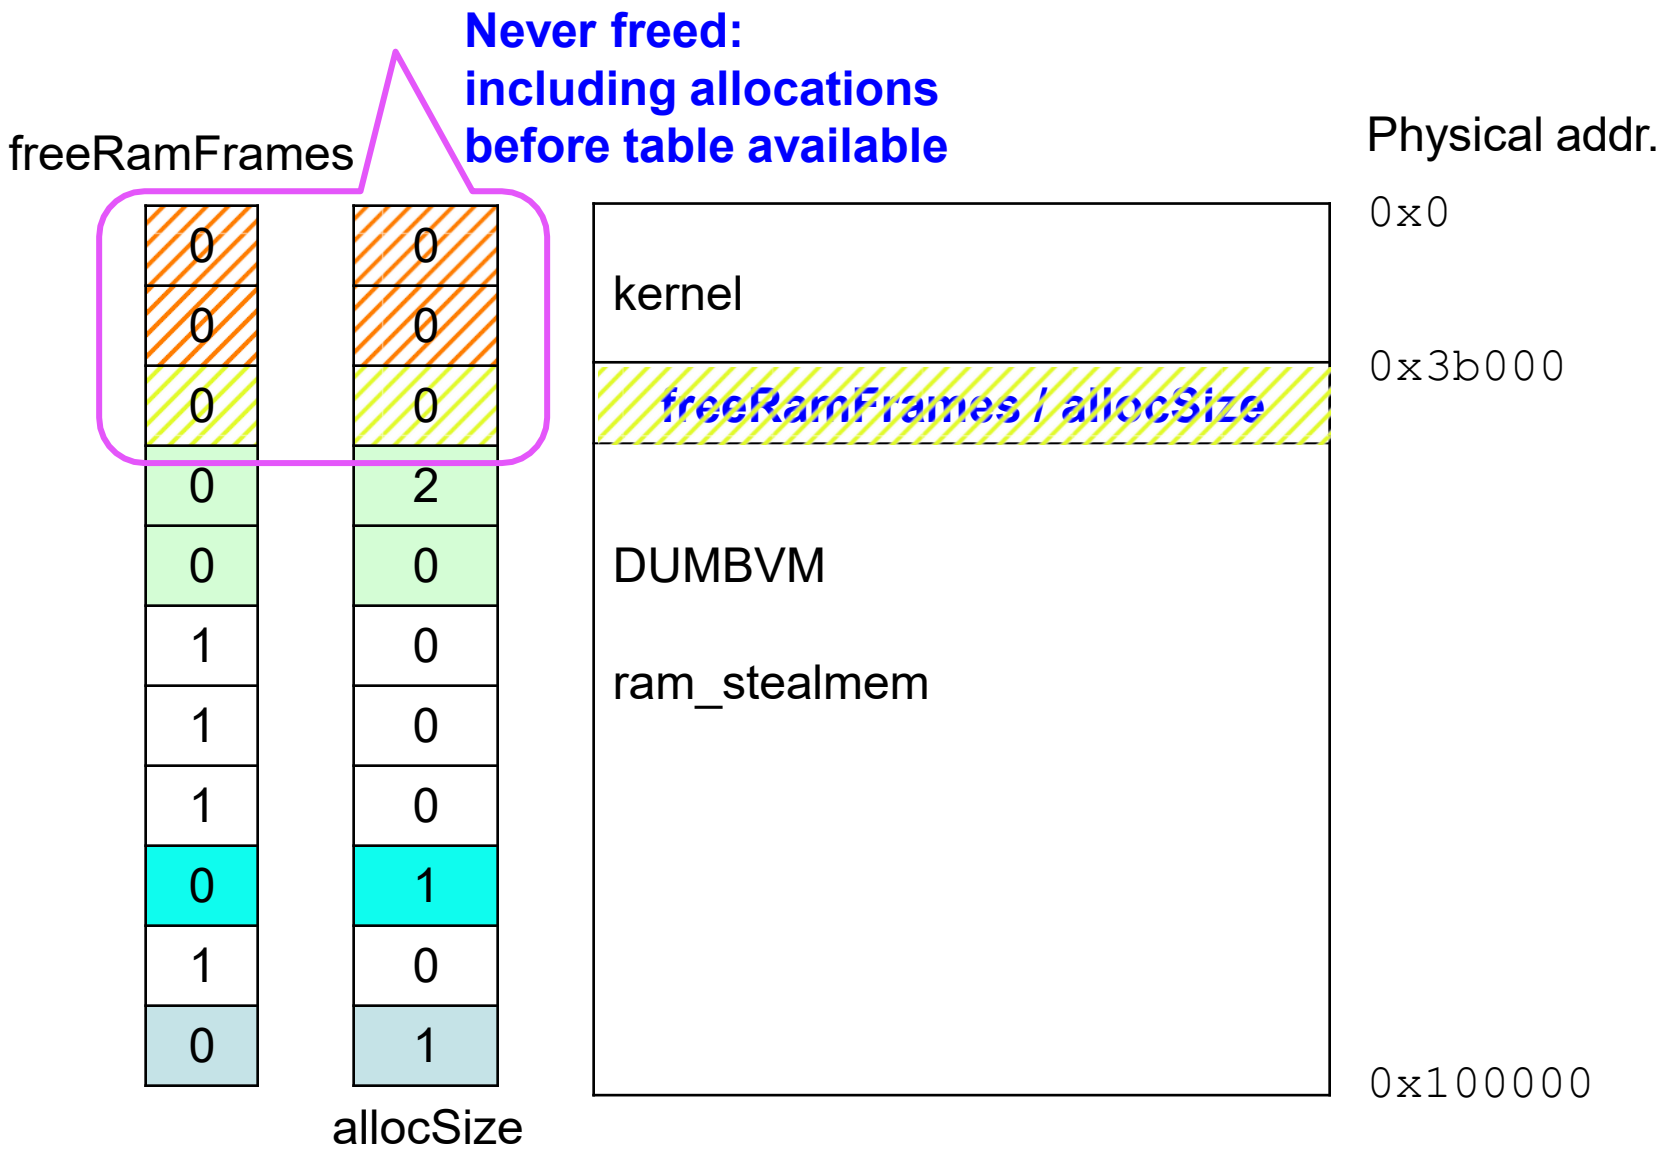
Physical addr.

0x0

0x3b000

0x100000

Dumbvm



Dumbvm

freeRamFrames

Physical addr.

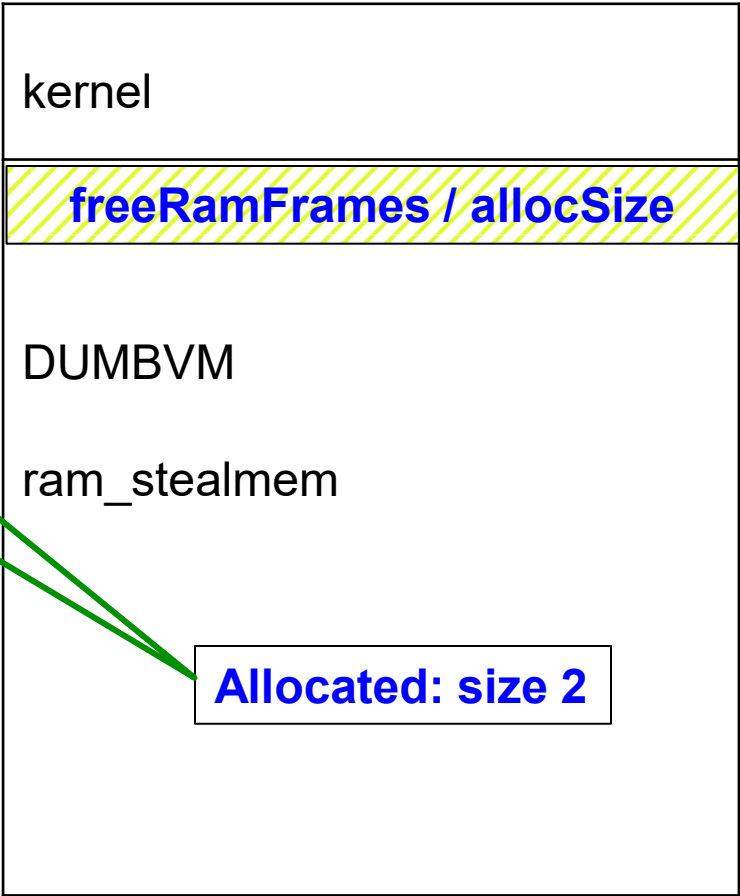
0x0

0x3b000

0x100000

0	0
0	0
0	0
0	2
0	0
1	0
1	0
1	0
0	1
1	0
0	1

allocSize



Allocated: size 2

Dumbvm

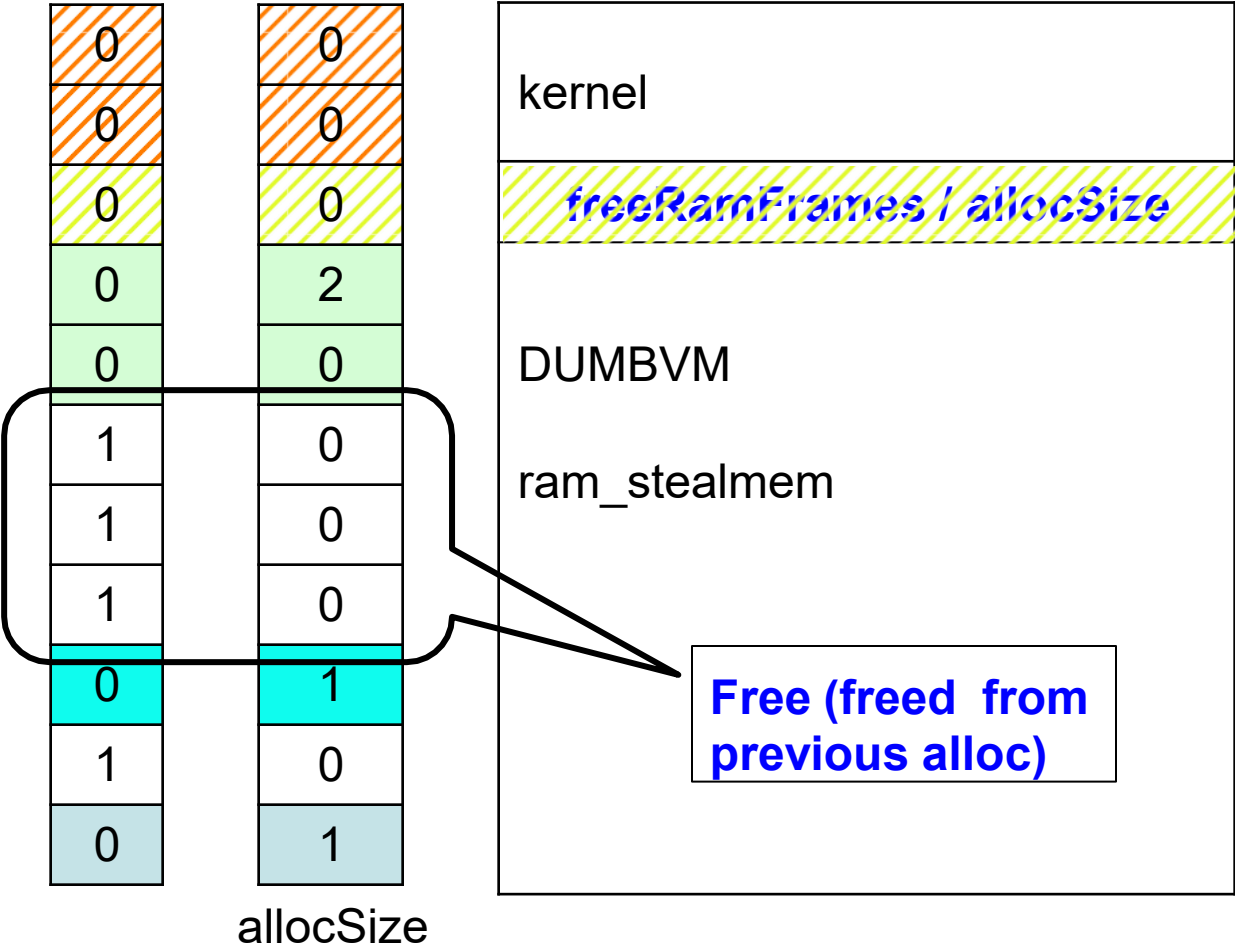
freeRamFrames

Physical addr.

0x0

0x3b000

0x100000



Free_kpages & as_destroy

```
void free_kpages(vaddr_t addr) {
    if (isTableActive()) {
        paddr_t paddr = addr - MIPS_KSEG0;
        long first = paddr/PAGE_SIZE;
        KASSERT(nRamFrames>first);
        freeppages(paddr, allocSize[first]);
    }
}

void as_destroy(struct addrspace *as) {
    dumbvm_can_sleep();
    freeppages(as->as_pbase1, as->as_npages1);
    freeppages(as->as_pbase2, as->as_npages2);
    freeppages(as->as_stackbase, DUMBVM_STACKPAGES);
    kfree(as);
}
```

Initialization

```
void vm_bootstrap(void) {
    int i;
    nRamFrames = ((int)ram_getsize())/PAGE_SIZE;
    /* alloc freeRamFrame and allocSize */
    freeRamFrames = kmalloc(sizeof(unsigned char)*nRamFrames);
    allocSize      = kmalloc(sizeof(unsigned long)*nRamFrames);
    if (freeRamFrames==NULL || allocSize==NULL) {
        /* reset to disable this vm management */
        freeRamFrames = allocSize = NULL; return;
    }
    for (i=0; i<nRamFrames; i++) {
        freeRamFrames[i] = (unsigned char)0; allocSize[i] = 0;
    }
    spinlock_acquire(&freemem_lock);
    allocTableActive = 1;
    spinlock_release(&freemem_lock);
}
```

getppages

```
static paddr_t getppages(unsigned long npages) {
    paddr_t addr;

    /* try freed pages first */
    addr = getfreepages(npages);
    if(addr== 0) { /* call stealmem */
        spinlock_acquire(&stealmem_lock);
        addr = ram_stealmem(npages);
        spinlock_release(&stealmem_lock);
    }
    if(addr != 0 && isTableActive()) {
        spinlock_acquire(&freemem_lock);
        allocSize[addr/PAGE_SIZE] = npages;
        spinlock_release(&freemem_lock);
    }
    return addr;
}
```


getfreeppages

```
static paddr_t getfreeppages(unsigned long npages) {
    paddr_t addr;
    long i, first, found, np = (long)npages;

    if (!isTableActive()) return 0;
    spinlock_acquire(&freemem_lock);
    // Linear search of free interval
    for (i=0, first=found=-1; i<nRamFrames; i++) {
        if (freeRamFrames[i]) {
            if (i==0 || !freeRamFrames[i-1])
                first = i; /* set first free in an interval */
            if (i-first+1 >= np)
                found = first;
        }
    }
}
```

getfreeppages

```
if (found>=0) {
    for (i=found; i<found+np; i++) {
        freeRamFrames[i] = (unsigned char)0;
    }
    allocSize[found] = np;
    addr = (paddr_t) found*PAGE_SIZE;
}
else {
    addr = 0;
}

spinlock_release(&freemem_lock);

return addr;
}
```

getfreepages

```
static int freepages(paddr_t addr, unsigned long npages){
    long i, first, np=(long)npages;
    if (!isTableActive()) return 0;
    first = addr/PAGE_SIZE;
    KASSERT(allocSize!=NULL);
    KASSERT(nRamFrames>first);

    spinlock_acquire(&freemem_lock);
    for (i=first; i<first+np; i++) {
        freeRamFrames[i] = (unsigned char)1;
    }
    spinlock_release(&freemem_lock);

    return 1;
}
```