

OS 161: Synchronization Primitives



Politecnico
di Torino

Department of Control and
Computer Engineering



System Programming - Sarah Azimi

CAD & Reliability Group
DAUIN- Politecnico di Torino

Chapter 6: Synchronization Tools

- Background
- The Critical-Section Problem
- Peterson's Solution
- Hardware Support for Synchronization
- Mutex Locks
- Semaphores
- Monitors
- Liveness
- Evaluation

Background

- Process can execute concurrently or in parallel.
- Coordinating access to shared resources
- Problems
 - Race conditions
 - Deadlocks
 - Resource starvation
- Solutions
 - Synchronization: locks, barriers, semaphores, etc.
 - ...

Critical Section Problem

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code that is shared with at least one other process
 - Process may be changing common variables, updating table, writing file, etc
 - When one process in critical section, no other may be in its critical section
- **Critical section problem** is to design protocol that process can use to synchronize their activity to cooperatively share data.
- The section of the code implementing this request is the **entry section**.
- The critical section may be followed by an **exit section**.
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, the remaining code is the **remainder section**

Critical Section

- General structure of process P_i

do {

entry section

critical section

exit section

remainder section

} while (true);

Solution to Critical-Section Problem

A solution to the critical-section problem must satisfy the following three requirements:

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.

Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non- preemptive

- **Preemptive** – allows a process to be preempted while it is running in kernel mode.
- **Non-preemptive** – does not allow a process running in kernel mode to be preempted, a kernel mode process will run until it exists kernel mode, blocks, or voluntarily yields control of CPU.
 - A non-preemptive kernel is essentially free of race conditions in kernel mode

Peterson's Solution

- Not guaranteed to work on modern architectures! (But good algorithmic description of solving the problem)
- It is restricted to two processes solution that alternate execution between their critical sections and remainder sections.
- Assume that the **load** and **store** machine-language instructions are atomic; that is, cannot be interrupted
- It requires the two processes to share two variables:
 - `int turn;`
 - `boolean flag[2]`
- The variable `turn` indicates whose turn it is to enter the critical section. If it is 1 then the process is allowed to execute in its critical section.
- The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = true` implies that process P_i is ready!

Algorithm for Process P_i

- To enter the critical section, the process p_i set the $flag[i]$ to be true and then sets $turn$ to the value j .
- Asserting that if the other process wishes to enter the critical section, it can do so.
- If both processes try to enter at the same time, $turn$ will be set to both i and j at roughly the same time.

```
while (true){  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /* remainder section */  
  
}
```

Peterson's Solution

- Provable that the three CS requirement are met:
 - Mutual exclusion is preserved
 - P_i enters CS only if:
 - either `flag[j] = false` or `turn = i`
 - Progress requirement is satisfied
 - Bounded-waiting requirement is met

Peterson's Solution

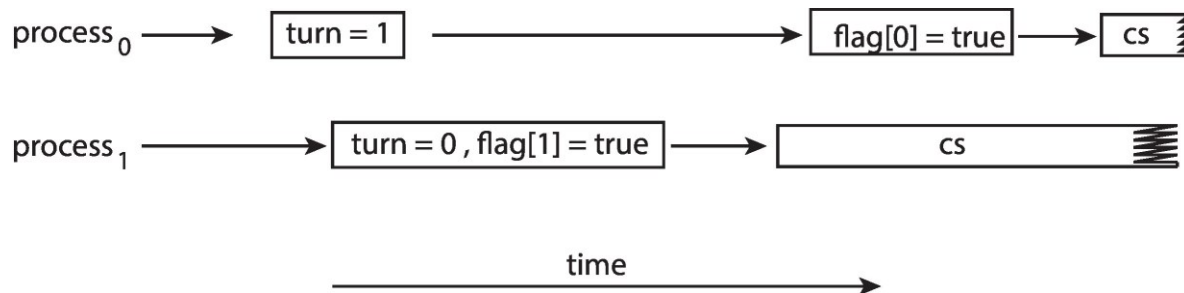
- Although useful for demonstrating an algorithm, Peterson's Solution is not guaranteed to work on modern architectures.
- To improve performance, processors and/or compilers may reorder read and write operations that have no dependencies.
- For single-threaded this is ok as the result will always be the same.
- For multithreaded the reordering may produce inconsistent or unexpected results!

Peterson's Solution

- Two threads share the data:
 - `boolean flag = false; int x = 0;`
- Thread 1 performs
 - `while (!flag)`
 - `;`
 - `print x`
- Thread 2 performs
 - `x = 100;`
 - `flag = true`
- What is the expected output?

Peterson's Solution

- 100 is the expected output.
- However, as there are no data dependencies between the variables flag and x, it is possible that a processor may reorder the instruction for thread 2 so that flag is assigned true before assignment of x=100.
 - `flag = true; x = 100;`
- If this occurs, the output may be 0!
- How this affect the Peterson's solution?



- This allows both processes to be in their critical section at the same time!
- The only way to preserve mutual exclusion is by using proper synchronization tools.

Synchronization Hardware

- Up to now, software-based solution to critical-section problem which do not guarantee to work on modern computer architecture.
- Three hardware instructions that provide support for solving the critical-section problem:
 - Memory barriers
 - Hardware instructions
 - Atomic variables

Memory Barriers

- **Hardware instructions to provide support for solving the critical-section problem:**
- **Memory model:** a computer architecture determines what memory guarantees it will provide to an application program.
- Memory models may be either:
 - **Strongly ordered** – where a memory modification of one processor is immediately visible to all other processors.
 - **Weakly ordered** – where a memory modification of one processor may not be immediately visible to all other processors.
- A **memory barrier** is an instruction that forces any change in memory to be propagated (made visible) to all other processors.

Memory Barriers

- We could add a memory barrier to the following instructions to ensure Thread 1 outputs 100:
- Thread 1 now performs

```
while (!flag)
    memory_barrier();
print x
```

- Thread 2 now performs

```
x = 100;
memory_barrier();
flag = true
```


Hardware Instructions

- Many modern computer systems provide special hardware instructions that allow us to either test and modify the content of a word or to swap the contents of two words **atomically**:
- **Test-and-Set** instruction >> performing a test on a condition, if the condition is true, set a value.
- **Compare-and-Swap** instruction >> comparing the two variables and swapping the values.

test_and_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = true; return rv;
}
```

1. Executed **atomically**
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to **true**

Solution using test_and_set()

- Shared boolean variable `lock`, initialized to `false`
- Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
  
    /* remainder section */  
} while (true);
```

compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
    return temp;  
}
```

1. Executed atomically
2. Returns the original value of passed parameter `value`
3. Set the variable `value` the value of the passed parameter `new_value` but only if `*value == expected` is true. That is, the swap takes place only under this condition.

Solution using compare_and_swap

- Shared integer `lock` initialized to 0;
- Solution:

```
while (true){  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
}
```

Atomic Variables

- Typically, instructions such as compare-and-swap are used as building blocks for other synchronization tools.
- One tool is an **atomic variable** that provides *atomic* (uninterruptible) operations on basic data types such as integers and booleans.
- For example, the `increment()` operation on the atomic variable
 - **sequence** ensures **sequence** is incremented without interruption:

Atomic Variables

- The functions are often implemented using `compare_and_swap()` operation:
- The `increment()` function can be implemented as follows:

```
void increment(atomic_int *v)
{
    int temp;

    do {
        temp = *v;
    }
    while (temp != (compare_and_swap(v, temp, temp+1)) );
}
```

Mutex Locks

- Previous solutions are complicated and generally inaccessible to application programmers
- OS designers build software tools to solve critical section problem.
- Simplest is **Mutex Lock** for protecting critical sections and prevent race condition:
 - A process must acquire a lock before entering a critical section using **acquire()** function
 - Releasing the lock when it exits the critical section using **release()** function
 - A mutex lock has a Boolean variable available whose value indicates if the lock is available or not.
 - If the lock is available, a call to **acquire()** succeeds and the lock is considered unavailable.

Mutex Locks

- Calls to `acquire()` and `release()` must be atomic.
 - Usually implemented via hardware atomic instructions such as compare-and-swap.
 - **Disadvantage:** this solution requires busy waiting, while the process is in critical section, the other processes must continuously loop to `acquire()`.
 - Wasting CPU cycle.
- This lock therefore called a **spinlock**.
 - **Advantage:** no context switch is required when a process might wait on a lock. If a lock is to be held for a short duration, one thread can “spin” on one processing core while another thread performs its critical section on one another.

Solution to Critical-section Problem Using Locks

```
while (true) {  
    acquire lock  
  
    critical section  
  
    release lock  
  
    remainder section  
}
```

OS/161 Spinlocks (kern/thread/spinlock.c)

```
spinlock_acquire(struct spinlock *splk) {
    ...
    while (1) {
        /* Do test-test-and-set, that is, read first before
           doing test-and-set, to reduce bus contention.
           Test-and-set is a machine-level atomic operation
           */
        if (spinlock_data_get(&splk->splk_lock) != 0) {
            continue;
        }
        if (spinlock_data_testandset(&splk->splk_lock) != 0) {
            continue;
        }
        break;
    }
    ...
}
```

Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- A semaphore **S** is an integer variable that, apart from initialization is accesses through two standard **atomic** operations:

- **wait()** and **signal()**

- Definition of the **wait()**

```
wait(S) {  
    while (S <= 0);  
    // busy wait  
    S--;  
}
```

- Definition of the **signal()**

```
signal(S) {  
    S++;  
}
```

Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
 - Behave similarly to **mutex lock**
- Can solve various synchronization problems
- Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “**synch**” initialized to 0

P1:

$S_1;$

signal(synch) ;

P2:

wait(synch); $S_2;$

- Can implement a counting semaphore **S** as a binary semaphore

Semaphore implementation

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time.
- Thus, the implementation becomes the critical section problem where the wait and signal code are placed in the critical section
- Could now have **busy waiting** in critical section implementation
 - But implementation code is short
 - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution

Semaphore implementation

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

OS/161: Disabling Interrupts

- On a uniprocessor, only one thread at a time is actually running.
- If the running thread is executing a critical section, mutual exclusion may be violated if:
 1. the running thread is preempted (or voluntarily yields) while it is in the critical section, and
 2. the scheduler chooses a different thread to run, and this new thread enters the same critical section that the preempted thread was in
- Since preemption is caused by timer interrupts, mutual exclusion can be enforced by disabling timer interrupts before a thread enters the critical section, and re-enabling them when the thread leaves the critical section. This is the way that the OS/161 kernel enforces mutual exclusion.

OS/161: Disabling Interrupts

- Since preemption is caused by timer interrupts, mutual exclusion can be enforced by disabling timer interrupts before a thread enters the critical section, and re-enabling them when the thread leaves the critical section. This is the way that the OS/161 kernel enforces mutual exclusion.
- There is a simple interface (`splhigh()`, `spl0()`, `splx()`) for disabling and enabling interrupts.
 - See `kern/arch/mips/include/spl.h`.
- In reality the change the priority of the threads. So, the threads with lower priority cannot interrupt the execution of a thread with higher priority.

OS/161 Semaphores (1.9x: interrupt based)

```
struct semaphore {  
    char *name; volatile  
    int count;  
};
```

Structure of data defining semaphore: name of the semaphore and count variable

```
struct semaphore *sem_create(const char *name, int  
    initial count);  
void P(struct semaphore *);  
void V(struct semaphore *);  
void sem_destroy(struct semaphore *);
```

see

- kern/include/synch.h
- kern/thread/synch.c

Four functions for creating, signaling, waiting and destroying semaphore.

OS/161 Semaphores: P() (1.9x: interrupt based)

```
void P(struct semaphore *sem) {
    int spl;
    assert(sem != NULL);

    /* May not block in an interrupt handler.
     * For robustness, always check, even if we can actually
     * complete the P without blocking. */
    assert(in_interrupt==0);

    spl = splhigh();
    while (sem->count==0) {
        thread_sleep(sem);
    }
    assert(sem->count>0);
    sem->count--;
    splx(spl);
}
```

Saving the current priority of the process to recover later and Increasing the priority of the process.

Considering the semaphore value, if the value of semaphore is 0, go to sleep.

If sem is bigger than 0, increasing sem and recovering the priority of the process.

Changing the priority to the original one.

OS/161 Semaphores: V() (1.9x: interrupt based)

```
void V(struct semaphore *sem)
{
    int spl;
    assert(sem != NULL);
    spl = splhigh();
    sem->count++;
    assert(sem->count>0);
    thread_wakeup(sem);
    splx(spl);
}
```

Saving the current priority of the process to recover later and Increasing the priority of the process.

Increasing the sem

Waking up one of the processors in the waiting list.

Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Thread Blocking in OS/161

Implementation of semaphores are based on these two functions.

- OS/161 thread library functions:
 - void `thread_sleep`(const void *addr)
blocks the calling thread on address addr
 - void `thread_wakeup`(const void *addr)
unblock threads that are sleeping on address addr
- `thread_sleep()` is much like `thread_yield()`. The calling thread voluntarily gives up the CPU, the scheduler chooses a new thread to run, and dispatches the new thread. However
 - after a `thread_yield()`, the calling thread is *ready* to run again as soon as it is chosen by the scheduler
 - after a `thread_sleep()`, the calling thread is blocked, and should not be scheduled to run again until after it has been explicitly unblocked by a call to `thread_wakeup()`.

OS/161 Locks

- OS/161 also uses a synchronization primitive called a *lock*. Locks are intended to be used to enforce mutual exclusion.

```
struct lock *mylock = lock_create("LockName");
```

```
lock_acquire(mylock);
```

```
    critical_section /* e.g., call to list_remove_head */  
lock_release(mylock);
```

- A ***lock is similar to a binary semaphore*** with an initial value of 1. However, locks also enforce an additional constraint: **the thread that releases a lock must be the same thread that most recently acquired it.**
- The system enforces this additional constraint to help ensure that locks are used as intended.

Problems with Semaphores

- Incorrect use of semaphore operations:
 - `signal (mutex) ... wait (mutex)`
 - `wait (mutex) ... wait (mutex)`
 - Omitting of `wait (mutex)` and/or `signal (mutex)`
- These – and others – are examples of what can occur when semaphores and other synchronization tools are used incorrectly.

Semaphore limitation

- Can be hard to reason about synchronization
- Reason for waiting is embedded in P() (wait()) operation
 - Wait on counter
 - If (count == 0) sleep
 - Other waiting conditions not possible
 - E.g. if ((x==0) && (y>0 || z>0)) sleep
- Wait on condition
 - Condition checked outside P()
 - BUT checking needs mutual exclusion (extra lock/mutex/semaphore)
 - When sleeping (as a result of checking condition), extra lock/mutex/ semaphore is OWNED
 - POSSIBLE DEADLOCK

Condition Variables

- While mutexes implement synchronization by controlling thread access to data, condition variables allow threads to synchronize based upon the actual value of data.
- In a critical section, a thread can suspend itself on a condition variable if the state of the computation is not right for it to proceed.
 - It will suspend by waiting on a condition variable
 - It will release the critical section lock (MUTEX)
 - When that condition variable is signaled, it will become ready again. It will attempt to reacquire that critical section lock and only then it will be able to proceed

Condition Variables

- Condition x, y:
- Two operations are allowed on a condition variables:
 - `x.wait()` – a process that invokes the operation is suspended until `x.signal()`
 - `x.signal()` – resumes one of the processes (if any) that invoked `x.wait()`
 - If no `x.wait()` on the variable, then it has no effect on the variable.

Condition Variables Choices

- If process P invokes **x.signal()** , and process Q is suspended in
 - **wait()** , what should happen next?
 - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
 - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
 - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
 - Both have pros and cons – language implementer can decide
 - Monitors implemented in Concurrent Pascal compromise
 - P executing signal immediately leaves the monitor, Q is resumed
 - Implemented in other languages including Mesa, C#, Java

Using Condition variables with cv_signal

- **Always** used together with locks
 - The lock protects the shared data that is modified and tested when deciding whether to wait or signal/broadcast
- General Usage:

P_i

```
lock_acquire(lock);
while (condition not true)
{
    cv_wait(cond, lock);
}
... // do stuff
lock_release(lock);
```

P_j

```
lock_acquire(lock);
... // modify condition
cv_signal(cond);
lock_release(lock);
```

Using Condition variables with cv_broadcast

P_i

```
lock_acquire(lock);
while (condition_i not true){
    cv_wait(cond, lock);
}
... // do stuff
lock_release(lock);
```

P_k

```
lock_acquire(lock);
while (condition_k not true){
    cv_wait(cond, lock);
}
... // do stuff
lock_release(lock);
```

P_j

```
lock_acquire(lock);
...
// modify conditions
// either for  $P_i$  or  $P_k$ 
cv_broadcast(cond);
lock_release(lock);
```

Sending the signal in broadcast at processes that are waiting, waking up all, the one that take the control of CPU first, goes on with the implementation. (to be implemented)

OS/161 Wait Channels

- Same as condition variables not using semaphores but using spinlocks (busy waiting).
 - Spinlock owned for short time
 - Nested or multiple spinlocks not allowed
- Kernel level synchronization objects
- Integrated with thread scheduling
 - Spinlock handled within `thread_switch`

OS/161 TODO

- Kernel level synchronization objects to be implemented:
 - Locks
 - Condition Variables
- Strategy:
 - look at semaphores
 - Use spinlocks and wait channels

OS/161 TODO

- Kernel level synchronization objects to be implemented:
 - Locks
 - Condition Variables
- Strategy:
 - look at semaphores
 - Use spinlocks and wait channels



LAB 3!