

File-System Implementation



Politecnico
di Torino

Department of Control and
Computer Engineering



System Programming - Sarah Azimi

CAD & Reliability Group
DAUIN- Politecnico di Torino

Introduction

- File system provides the mechanism for on-line storage and access to file contents, including data and programs.
- File systems usually reside permanently on secondary storage, which is designed to hold a large amount of data.
- This chapter is concerned with issues surrounding file storage and access on the most common secondary-storage media, hard disk drives and nonvolatile memory devices.

Objective

- Describe the details of implementing local file systems and director structures
- Discuss block allocation and free-block algorithms and trade-offs
- Explore file system efficiency and performance issues
- Look at recovery from file system failures describe the WAFL file system as an example

File System Structure


- Disks provide most of the secondary storage on which file systems are maintained.
 - A disk can be rewritten in place. It is possible to read a block from the disk, modify the block, and write it back into the same block.
 - A disk can access directly any block of information it contains.

File System Structure

- File systems provide efficient and convenient access to the storage device by allowing data to be stored, located and retrieved easily.
- A file system has two different design problems:
 - Defining how the file system should look to the user
 - Defining a file and its attributes, the operation allowed on a file, the directory structure for organizing files.
 - Creating algorithms and data structures to map the logical file system onto the physical secondary-storage devices.


Layered File System

- File system is composed of many different levels:
 - I/O control
 - Basic file system
 - File-organization module
 - Logical file system
 - Application programs




Layered File System

- File system is composed of many different levels:
 - I/O control
 - **Device drivers** and interrupt handlers to transfer information between the main memory and the disk system.
 - Device driver can be thought of a translator.
 - Basic file system
 - Issuing generic commands to the appropriate device driver to read and write blocks on the storage device.
 - Managing the memory buffers and caches holding various file system, directory and data blocks.
 - File-organization module
 - Knows about files and their logical blocks.
 - Each file's logical blocks are numbered from 0 (or 1) through N.
 - Including the free-space manager for tracking unallocated blocks and providing these blocks to the file-organization module when requested.



Layered File System

- File system is composed of many different levels:
 - I/O control
 - Basic file system
 - File-organization module
 - Logical file system
 - Managing **metadata** information including all of the file-system structure except the contents of the files.
 - Managing the directory structure to provide the file-organization module with the information needed.
 - Maintaining file structure via **file-control blocks** containing information about the file, such as ownership, permissions and location of the file contents.
 - Application programs



Layered File System

- Benefits of layer structure:
 - Duplication of code is minimized. The I/O control and sometimes the basic file-system code can be used by multiple file systems.
 - Each file system can have its own logical file-system and file-organization modules.
- Disadvantages:
 - Layering can introduce more operating-system overhead, which may result in decreased performance

File System Layers

- Many file systems, sometimes many within an operating system
 - Each with its own format (CD-ROM is ISO 9660; Unix has **UFS**, FFS; Windows has FAT, FAT32, NTFS as well as floppy, CD, DVD Blu-ray, Linux has more than 40 types, with **extended file system** ext2 and ext3 leading; plus distributed file systems, etc.)
 - New ones still arriving – ZFS, GoogleFS, Oracle ASM, FUSE

File System Implementation

- We have system calls at the API level, but how do we implement their functions?
 - On-disk and in-memory structures
- **Boot control block** contains info needed by system to boot OS from that volume
 - If the disc does not contain an operating system, this block can be empty
 - Needed if volume contains OS, usually first block of volume
- **Volume control block** (superblock, master file table) contains volume details
 - Total number of blocks, number of free blocks, block size, free block pointers or array
- Directory structure organizes the files
 - Names and inode numbers, stored in master file table

File System Implementation

- Per-file **File Control Block (FCB)** contains many details about the file such as ownership, permissions and location of the file contents:
 - It has a unique identifier number to allow association with a directory entry.
 - inode number, permissions, size, dates
 - NFTS stores into in master file table using relational DB structures

In-Memory File System Structures

- The in-memory information is used for both file-system management and performance improvement via caching.
- The data are loaded at the mount time, updated during file-system operation and discarded at dismount.
- Several types structures may be included:
 - An in-memory **mount table** contains information about each mounted volume.
 - An in-memory directory structure cache holds the directory information of recently accessed directories.
 - The **system-wide open-file table** contains a copy of the FCB of each open file, as well as other information.
 - Buffers hold file-system blocks when they are being read from or written to a file system.

In-Memory File System Structures

- To create a new file system, a process calls the logical file system.
 - The logical file system knows the format of the directory structures.
 - To create a new file, it allocates a new FCB.
 - The system then read the appropriate directory into memory, updates it with the new file name and FCB and writes it back to the file system.

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

Partitions and Mounting

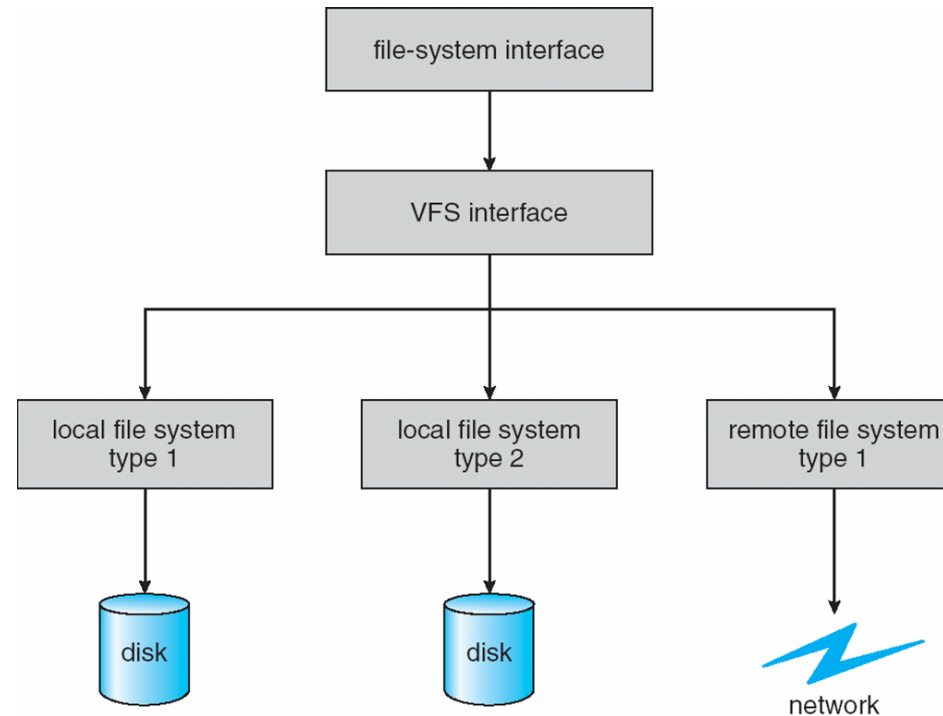
- Partition can be a volume containing a file system (“cooked”) or **raw** – just a sequence of blocks with no file system
- Boot block can point to boot volume or boot loader set of blocks that contain enough code to know how to load the kernel from the file system
 - Or a boot management program for multi-os booting
- **Root partition** contains the OS, other partitions can hold other Oses, other file systems, or be raw
 - Mounted at boot time
 - Other partitions can mount automatically or manually
- At mount time, file system consistency checked
 - Is all metadata correct?
 - If not, fix it, try again
 - If yes, add to mount table, allow access

Virtual File Systems

- **Virtual File Systems (VFS)** on Unix provide an object-oriented way of implementing file systems
- VFS allows the same system call interface (the API) to be used for different types of file systems
 - Separates file-system generic operations from implementation details
 - Implementation can be one of many file systems types, or network file system
 - Implements **vnodes** which hold inodes or network file details
 - Then dispatches operation to appropriate file system implementation routines

Virtual File Systems

- The API is to the VFS interface, rather than any specific type of file system



Virtual File System Implementation

- For example, Linux has four object types:
 - inode, file, superblock, dentry
- VFS defines set of operations on the objects that must be implemented
 - Every object has a pointer to a function table
 - Function table has addresses of routines to implement that function on that object
 - For example:
 - `•int open(. . .)` —Open a file
 - `•int close(. . .)` —Close an already-open file
 - `•ssize_t read(. . .)` —Read from a file
 - `•ssize_t write(. . .)` —Write to a file
 - `•int mmap(. . .)` —Memory-map a file

Directory Implementation

- Directory allocation and directory management algorithms significantly affects the efficiency, performance and reliability of the file system.
 - Linear List
 - Hash Table

Directory Implementation

- Directory allocation and directory management algorithms significantly affects the efficiency, performance and reliability of the file system.
 - Linear List
 - Using the linear list of file names with pointers to the data blocks.
 - This method is simple to program but time-consuming to execute.
 - To create a new file, first searching the directory to be sure that no existing file has the same name.
 - Then, adding a new entry at the end of the directory.
 - To delete a file, searching the directory for the named file and then releasing the space allocated to it.
 - The disadvantage of a linear list of directory is that finding a file requires a linear search.
 - Directory information is used frequently and users will notice if access to it is slow.

Directory Implementation

- Directory allocation and directory management algorithms significantly affects the efficiency, performance and reliability of the file system.
 - Hash Table
 - A linear list stores the directory, but a hash data structure is also used.
 - The hash table takes a value computed from the file name and returns a pointer to the file name in the linear list.
 - It greatly decrease the directory search time.
 - The major difficulties with a hash table are its generally fixed size and the dependence of the hash function on that size.

Allocation Methods


- An allocation method refers to how disk blocks are allocated for files:
 - Contiguous allocation
 - Linked allocation
 - Indexed allocation

Allocation Methods - Contiguous


- An allocation method refers to how disk blocks are allocated for files:
- **Contiguous allocation** – each file occupies set of contiguous blocks
 - Best performance in most cases
 - Simple – only starting location (block #) and length (number of blocks) are required
 - Accessing file:
 - For sequential access, the file system remembers the address of the last block references and read the next block.
 - For direct access to block i of a file that starts at block b , we can immediately access block $b+i$.

Contiguous Allocation

- Mapping from logical to physical



Block to be accessed = Q +
starting address
Displacement into block = R



directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

Extent-Based Systems

- **Problems of contiguous allocation:**
- Suffering from external fragmentation.
 - Files are allocated and deleted; the free storage space is broken into little pieces.
 - Largest amount of contiguous memory is not sufficient for a request.
 - **Solution:** copy an entire file system onto another device and copying the file back onto the original device by allocating contiguous space from this one large hole: **compacting** all free space into one large hole.

Extent-Based Systems

- **Problems of contiguous allocation:**
- Determining how much space is needed for a file.
 - When a file is created, the total amount of space it will need must be found and allocated which is difficult to estimate.
 - If we allocate too little space to a file, we may find the file cannot be extended.

Solutions:

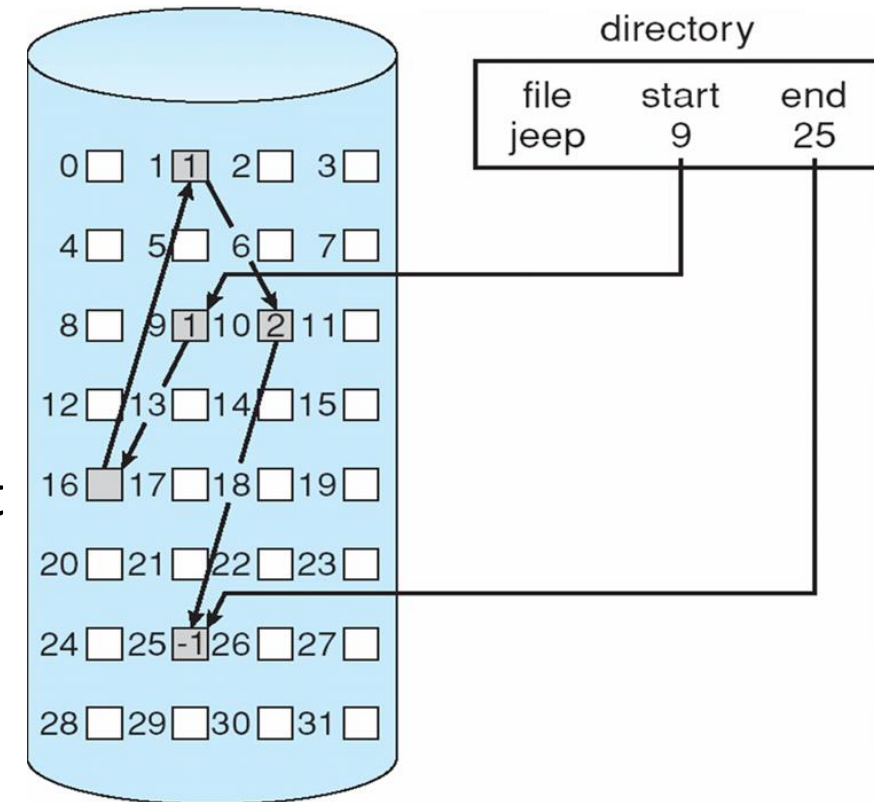
1. The user program can be terminated with an error message and run the program again.
 2. Finding a larger hole, copying the contents of the file to the new space and release the previous space.
- Time Consuming Solutions

Extent-Based Systems

- **Problems of contiguous allocation:**
- Determining how much space is needed for a file.
 - Modified contiguous-allocation scheme:
 - A contiguous chunk of space is allocated initially.
 - If the amount proves not to be large enough, another chunk of contiguous space, known as an **extent** is added.
 - The location of a file's blocks is then recorded as a location and a block count, plus a link to the first block of the next extent.


Allocation Methods - Linked

- **Linked allocation** – each file a linked list of storage blocks
 - File ends at nil pointer
 - No external fragmentation
 - **Each block contains pointer to next block**
 - No compaction, external fragmentation
 - Free space management system called when new block needed
 - Improve efficiency by clustering blocks into groups but increases internal fragmentation
 - Reliability can be a problem
 - Locating a block can take many I/Os and disk seeks
 - Disadvantages: it can be used only for sequential access files. To find the i^{th} block of file, we must start from the beginning of the file.




Allocation Methods - Linked

- An important variation on linked allocation is the use of a **File Allocation Table (FAT)** variation
- A section of storage at the beginning of each volume is set aside to contain the table.
- The table has one entry for each block and is indexed by block number.
- The directory entry contains the block number of the first block of the files.
- The table entry indexed by that block number contains the block number of the next block in the file.
- This chain continues until it reached the last block which has a special end-of-file value as the table entry.




File Allocation Table

- An example of FAT structure for a file consisting of disk blocks 217, 618, and 339.



Allocation Methods - Indexed


- Linked allocation solves the external-fragmentation and size-declaration problems of contiguous allocation.
- In the absence of a FAT, linked allocation cannot support efficient direct access.
- **Indexed allocation** solve this problem by bringing all the pointers together into one location: the **index block**
 - Each file has its own index block(s) of pointers to its data blocks



index table


Example of Indexed Allocation

- Each file has its own index block which is an array of storage-block addresses.
- The i^{th} entry in the index block points to the i^{th} block of the file.
- The directory contains the address of the index block.
- To find and read the i^{th} block, we use the pointer in the i^{th} index-block entry.



Indexed Allocation

- Need index table
- Random access
- Dynamic access without external fragmentation, but have overhead of index block
- Mapping from logical to physical in a file of maximum size of 256K bytes and block size of 512 bytes. We need only 1 block for index table



Q = displacement into index table

R = displacement into block

Indexed Allocation - Mapping

- Mapping from logical to physical in a file of unbounded length (block size of 512 words)
- Linked scheme – Link blocks of index table (no limit on size)

$$LA / (512 \times 511) \begin{cases} Q_1 \\ R_1 \end{cases}$$

Q_1 = block of index table

R_1 is used as follows:

$$R_1 / 512 \begin{cases} Q_2 \\ R_2 \end{cases}$$

Q_2 = displacement into block of index table

R_2 displacement into block of file:

Indexed Allocation - Mapping

- Two-level index (4K blocks could store 1,024 four-byte pointers in outer index - > 1,048,567 data blocks and file size of up to 4GB)


$$LA / (512 \times 512) \begin{cases} Q_1 \\ R_1 \end{cases}$$

Q_1 = displacement into outer-index
 R_1 is used as follows:

$$R_1 / 512 \begin{cases} Q_2 \\ R_2 \end{cases}$$

Q_2 = displacement into block of index table
 R_2 displacement into block of file:

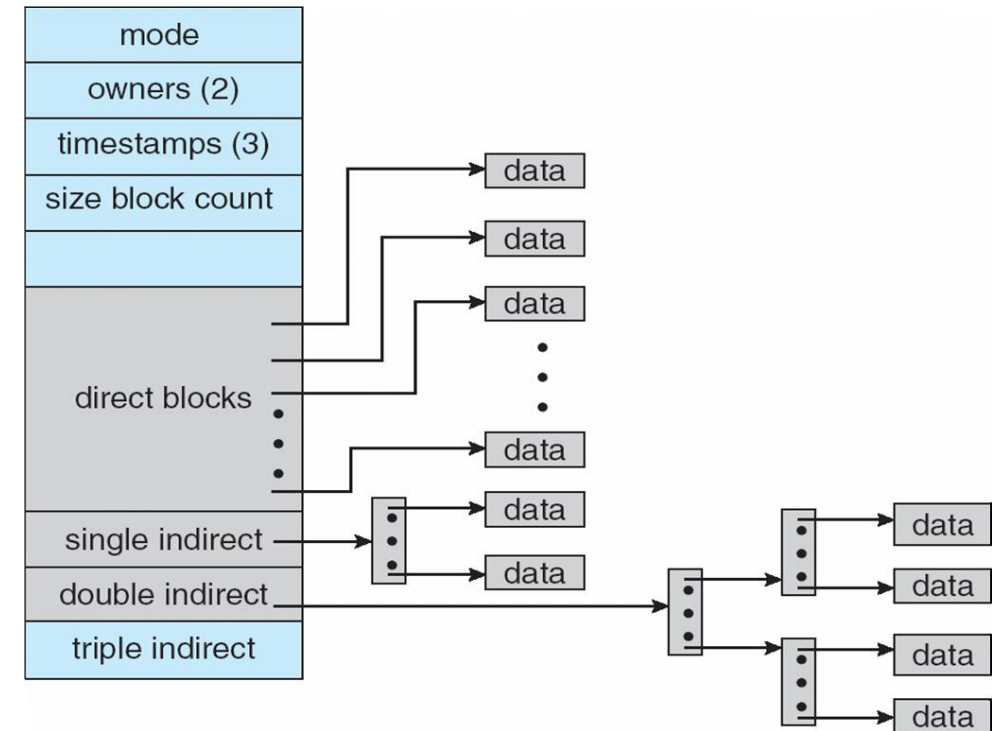
Indexed Allocation - Mapping



Combined Scheme: UNIX UFS

- TBD

4K bytes per block, 32-bit addresses



More index blocks than can be addressed with 32-bit file pointer

Performance

- The allocation methods vary in their storage efficiency and data-block access time.
- Before selecting an allocation method, we need to determine how the system will be used.
 - The system with mostly sequential access should not use the same method as a system with mostly random access.

Performance

- The allocation methods vary in their storage efficiency and data-block access time.
- Before selecting an allocation method, we need to determine how the system will be used.
 - **Contiguous allocation requires** only one access to get a block.
 - We can easily keep the initial address of the file in memory, we can calculate the address of the i^{th} block.
 - For linked allocation keep the address of the next block in memory and read it directly.

Performance

- The allocation methods vary in their storage efficiency and data-block access time.
- Before selecting an allocation method, we need to determine how the system will be used.
 - For **linked allocation** keep the address of the next block in memory and read it directly.
 - For sequential access it works, for direct access, an access to the i^{th} block might require i block read.

Performance

- The allocation methods vary in their storage efficiency and data-block access time.
- Before selecting an allocation method, we need to determine how the system will be used.
 - In **indexed allocation**, if the index block is already in memory, then the access can be made directly.
 - Keeping the index block in memory requires considerable space.
 - If this memory space is not available, then we should read first the index block and then the desired data block.

Performance

- The allocation methods vary in their storage efficiency and data-block access time.
- Before selecting an allocation method, we need to determine how the system will be used.
 - Therefore, some systems support direct-access files by using contiguous allocation and sequential-files by using linked allocation.
 - The types of access must be declared when the file is created.
 - Some systems combine contiguous allocation with indexed allocation:
 - Using contiguous allocation for small files
 - Automatically switching to an indexed allocation if the file grows large.

Performance

- Adding instructions to the execution path to save one disk I/O is reasonable
 - Intel Core i7 Extreme Edition 990x (2011) at 3.46Ghz = 159,000 MIPS
 - http://en.wikipedia.org/wiki/Instructions_per_second
 - Typical disk drive at 250 I/Os per second
 - $159,000 \text{ MIPS} / 250 = 630$ million instructions during one disk I/O
 - Fast SSD drives provide 60,000 IOPS
 - $159,000 \text{ MIPS} / 60,000 = 2.65$ millions instructions during one disk I/O

Free-space management


- Since storage space is limited, we need to reuse the space from deleted files for new files.
- File system maintains **free-space list** to track available blocks/clusters.
- To create a file system:
 - Search free space list for the required amount of space and allocate that space to the new file.
 - The space is then removed from the free-space list.
 - When the file is deleted, its space is added to the free-space list.

Free-space management

- The free-space list is implemented as **bitmap** or **bit vector**.
- Each block is represented by 1 bit.
 - If the block is free, the bit is 1. if the block is allocated, the bit is 0.
- For example, consider a disk where blocks 2,3,4,5,8,9,10,11,12,13,17,18,25,26, and 27 are free and the rest of the blocks are allocated. The free space bitmap would be:
 - 001111001111110001100000011100000...
- Advantages: Simplicity and efficiency in finding the first free block

Free-space management

- One technique for finding the first free block is to sequentially check each word in bitmap to see whether that value is not 0.
 - A 0 valued-word contains only 0 bits and represents a set of allocated blocks.
 - The first non-0 word is scanned for the first 1 bit.



$$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{block}[i] \text{ free} \\ 0 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

Block number calculation

(number of bits per word) * (number of 0-value words) + offset of first 1 bit

CPUs have instructions to return offset within word of first “1” bit

Free-space management

- Benefits: its relative simplicity and its efficiency in finding the first free block or n consecutive free blocks on disc.
- Drawbacks: inefficient unless the entire vector is kept in main memory. Keeping it in memory is possible for smaller devices but not necessarily for larger ones.
 - A 1.3 GB disk with 512-byte blocks would need a bitmap over 332KB to track its free blocks, although clustering blocks in groups of four reduces this number to around 83KB per disk.
 - A 1TB disk with 4KB blocks would require 32MB to store its bitmap. ($2^{40}/2^{12} = 2^{28}$ bits = 2^{25} bytes = 2⁵MB)

Free-space management

- Bit map requires extra space

- Example:

block size = 4KB = 2^{12} bytes

disk size = 2^{40} bytes (1 terabyte)


$n = 2^{40}/2^{12} = 2^{28}$ bits (or 32MB)

if clusters of 4 blocks -> 8MB of memory

- Easy to get contiguous files

Linked Free Space List on Disk

- Another approach to free-space management is to link together all the free blocks, keeping a pointer to the first free block in a special location in file system and caching it in memory.
- The first block contains a pointer to the next free block and so on.
- Not efficient since to traverse the list, we must read each block , which requires substantial I/O times on HDDs.



Free Space Management

- Grouping
 - Modify linked list to store address of next $n-1$ free blocks in first free block, plus a pointer to next block that contains free-block-pointers (like this one)
- Counting
 - Because space is frequently contiguously used and freed, with contiguous-allocation allocation, extents, or clustering
 - Keep address of first free block and count of following free blocks
 - Free space list then has entries containing addresses and counts

Free Space Management

- Space Maps
 - Used in **ZFS**
 - Consider meta-data I/O on very large file systems
 - Full data structures like bit maps couldn't fit in memory -> thousands of I/Os
 - Divides device space into **metaslab** units and manages metaslabs
 - Given volume can contain hundreds of metaslabs
 - Each metaslab has associated space map
 - Uses counting algorithm
 - But records to log file rather than file system
 - Log of all block activity, in time order, in counting format
 - Metaslab activity -> load space map into memory in balanced-tree structure, indexed by offset
 - Replay log into that structure
 - Combine contiguous free blocks into single entry

Efficiency and Performance

- Efficiency depends on:
 - Disk allocation and directory algorithms
 - Types of data kept in file's directory entry
 - Pre-allocation or as-needed allocation of metadata structures
 - Fixed-size or varying-size data structures

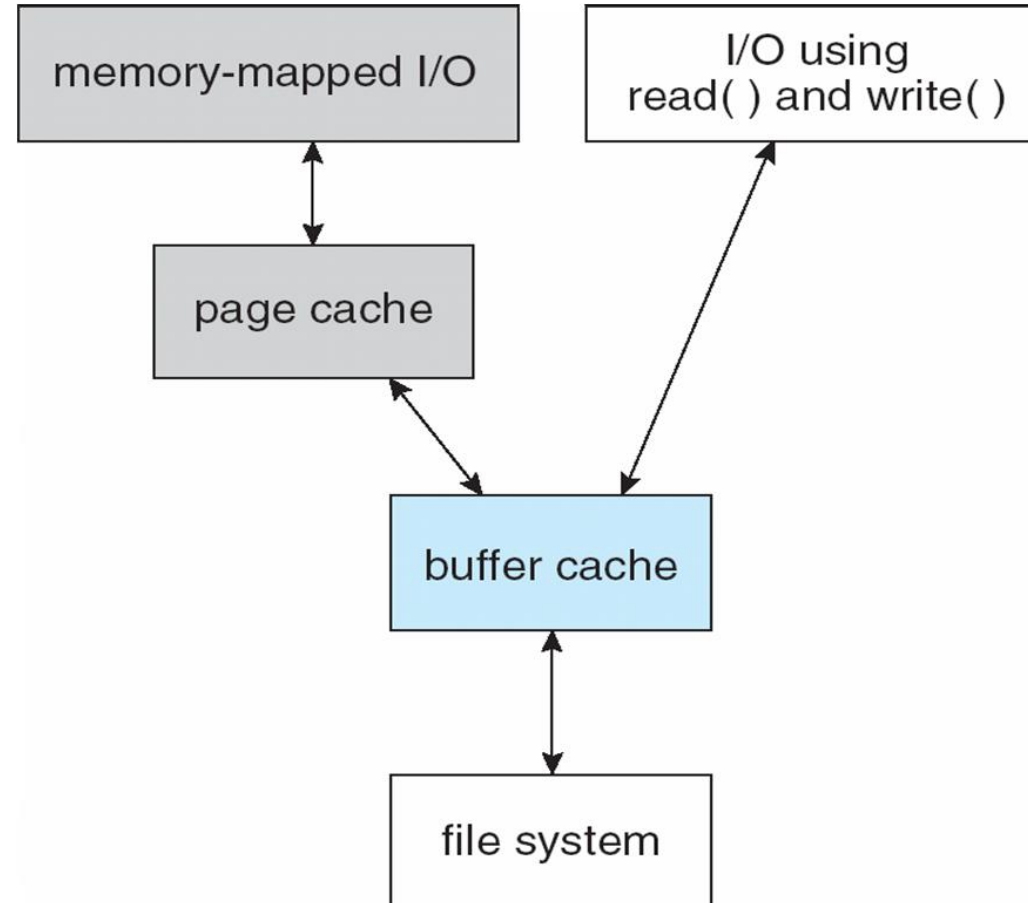
Efficiency and Performance

- Performance
 - Keeping data and metadata close together
 - **Buffer cache** – separate section of main memory for frequently used blocks
 - **Synchronous** writes sometimes requested by apps or needed by OS
 - No buffering / caching – writes must hit disk before acknowledgement
 - **Asynchronous** writes more common, buffer-able, faster
 - **Free-behind** and **read-ahead** – techniques to optimize sequential access
 - Reads frequently slower than writes

Page Cache

- A **page cache** caches pages rather than disk blocks using virtual memory techniques and addresses
- Memory-mapped I/O uses a page cache
- Routine I/O through the file system uses the buffer (disk) cache
- This leads to the following figure


I/O Without a Unified Buffer Cache



Unified Buffer Cache

- A **unified buffer cache** uses the same page cache to cache both memory-mapped pages and ordinary file system I/O to avoid **double caching**
- But which caches get priority, and what replacement algorithms to use?

I/O Using a Unified Buffer Cache



Recovery

- Files and directories are kept in main memory and on the storage volume
- Care must be taken to ensure that a system failure does not result in loss of data or in data inconsistency.
 - **Consistency checking** – compares data in directory structure with data blocks on disk, and tries to fix inconsistencies
 - Can be slow and sometimes fails
- Use system programs to **back up** data from disk to another storage device (magnetic tape, other magnetic disk, optical)
- Recover lost file or disk by **restoring** data from backup

Log Structured File Systems

- **Log structured** (or **journaling**) file systems record each metadata update to the file system as a **transaction**
- All transactions are written to a log
 - A transaction is considered committed once it is written to the log (sequentially)
 - Sometimes to a separate device or section of disk
 - However, the file system may not yet be updated
- The transactions in the log are asynchronously written to the file system structures
 - When the file system structures are modified, the transaction is removed from the log
- If the file system crashes, all remaining transactions in the log must still be performed
- Faster recovery from crash, removes chance of inconsistency of metadata

The Sun Network File System (NFS)

- An implementation and a specification of a software system for accessing remote files across LANs (or WANs)
- The implementation is part of the Solaris and SunOS operating systems running on Sun workstations using an unreliable datagram protocol (UDP/IP protocol and Ethernet)


The Sun Network File System (NFS)

- Interconnected workstations viewed as a set of independent machines with independent file systems, which allows sharing among these file systems in a transparent manner
 - A remote directory is mounted over a local file system directory
 - The mounted directory looks like an integral subtree of the local file system, replacing the subtree descending from the local directory
 - Specification of the remote directory for the mount operation is nontransparent; the host name of the remote directory has to be provided
 - Files in the remote directory can then be accessed in a transparent manner
 - Subject to access-rights accreditation, potentially any file system (or directory within a file system), can be mounted remotely on top of any local directory


The Sun Network File System (NFS)

- NFS is designed to operate in a heterogeneous environment of different machines, operating systems, and network architectures; the NFS specifications independent of these media
- This independence is achieved through the use of RPC primitives built on top of an External Data Representation (XDR) protocol used between two implementation-independent interfaces
- The NFS specification distinguishes between the services provided by a mount mechanism and the actual remote-file-access services

Three Independent File Systems




Mounting in NFS



(a)

Mounts



(b)

Cascading mounts

NFS Mount Protocol

- Establishes initial logical connection between server and client
- Mount operation includes name of remote directory to be mounted and name of server machine storing it
 - Mount request is mapped to corresponding RPC and forwarded to mount server running on server machine
 - Export list – specifies local file systems that server exports for mounting, along with names of machines that are permitted to mount them
- Following a mount request that conforms to its export list, the server returns a file handle—a key for further accesses
- File handle – a file-system identifier, and an inode number to identify the mounted directory within the exported file system
- The mount operation changes only the user's view and does not affect the server side

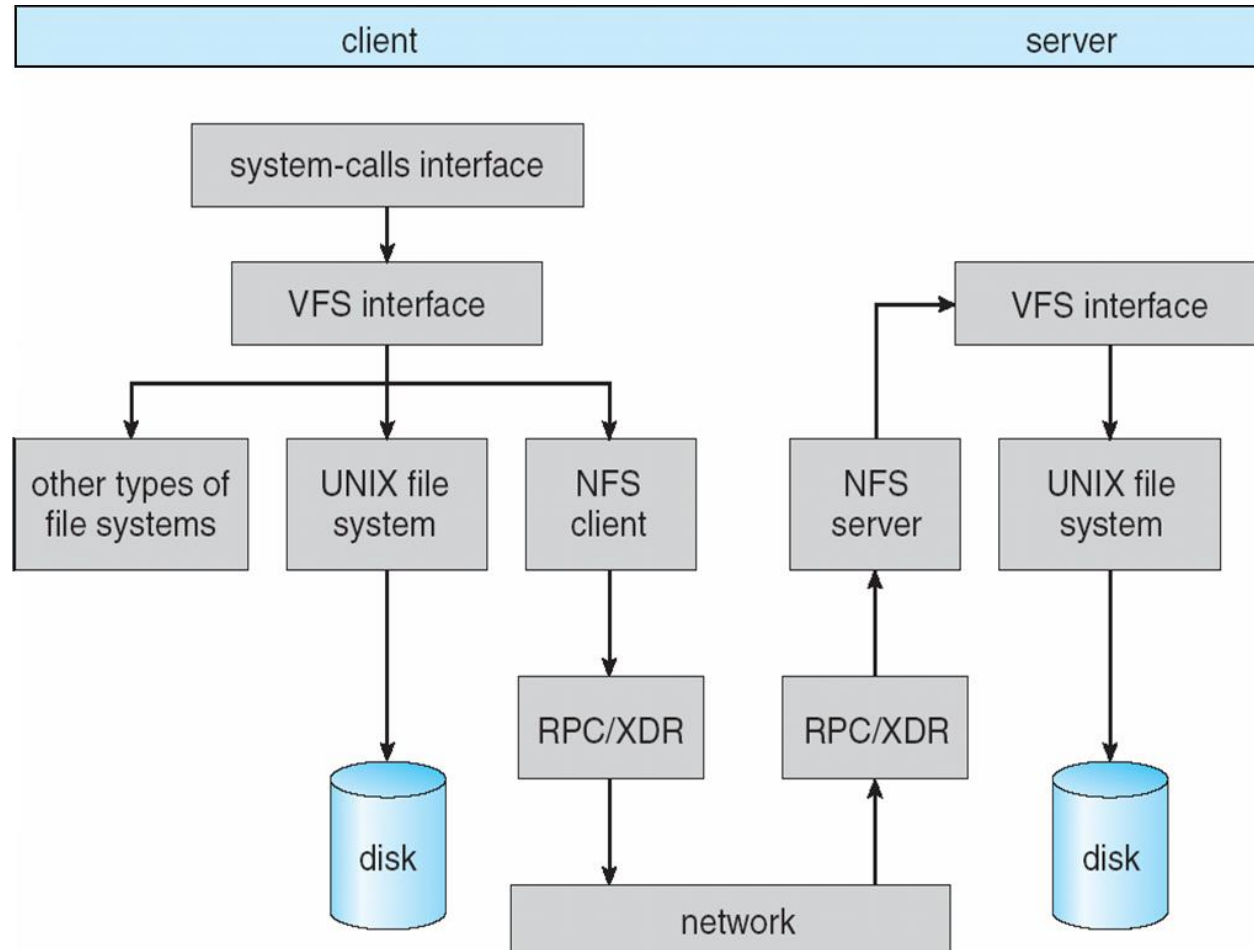
NFS Mount Protocol

- Provides a set of remote procedure calls for remote file operations. The procedures support the following operations:
 - searching for a file within a directory
 - reading a set of directory entries
 - manipulating links and directories
 - accessing file attributes
 - reading and writing files
- NFS servers are **stateless**; each request has to provide a full set of arguments (NFS V4 is just coming available – very different, stateful)
- Modified data must be committed to the server's disk before results are returned to the client (lose advantages of caching)
- The NFS protocol does not provide concurrency-control mechanisms

Three Major Layers of NFS Architecture

- UNIX file-system interface (based on the **open**, **read**, **write**, and **close** calls, and **file descriptors**)
- Virtual File System (VFS) layer – distinguishes local files from remote ones, and local files are further distinguished according to their file-system types
 - The VFS activates file-system-specific operations to handle local requests according to their file-system types
 - Calls the NFS protocol procedures for remote requests
- NFS service layer – bottom layer of the architecture
 - Implements the NFS protocol

Schematic View of NFS Architecture



NFS Path-Name Translation

- Performed by breaking the path into component names and performing a separate NFS lookup call for every pair of component name and directory vnode
- To make lookup faster, a directory name lookup cache on the client's side holds the vnodes for remote directory names

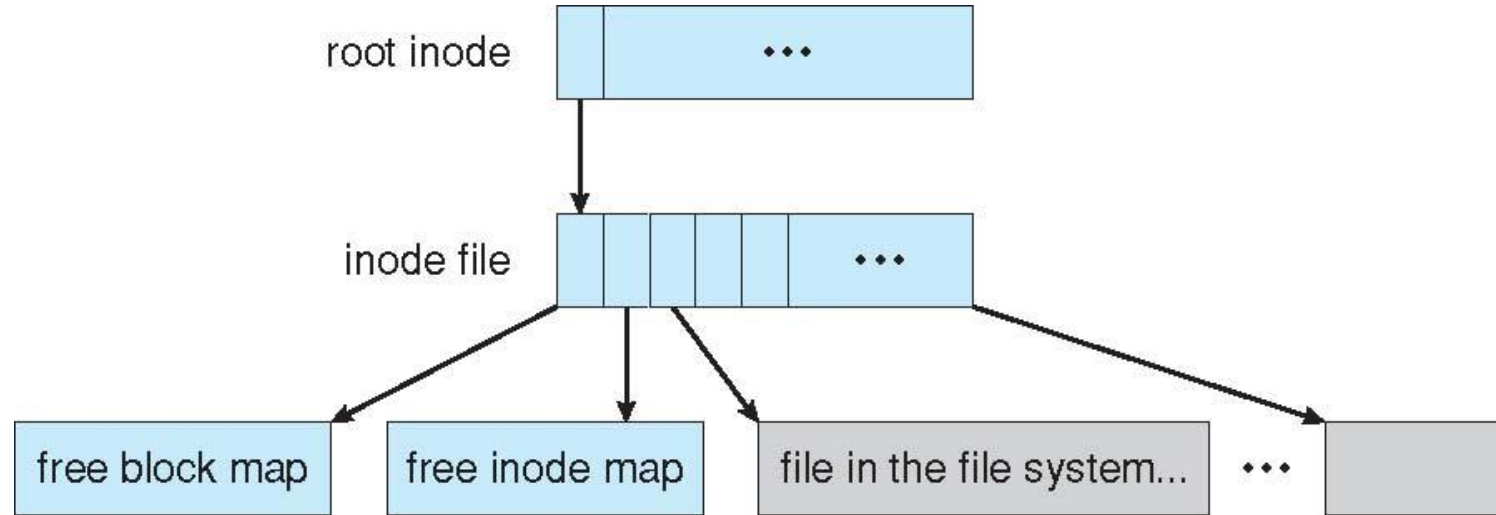
NFS Remote Operations

- Nearly one-to-one correspondence between regular UNIX system calls and the NFS protocol RPCs (except opening and closing files)
- NFS adheres to the remote-service paradigm, but employs buffering and caching techniques for the sake of performance
- File-blocks cache – when a file is opened, the kernel checks with the remote server whether to fetch or revalidate the cached attributes
 - Cached file blocks are used only if the corresponding cached attributes are up to date
- File-attribute cache – the attribute cache is updated whenever new attributes arrive from the server
- Clients do not free delayed-write blocks until the server confirms that the data have been written to disk


Example: WAFL File System

- Used on Network Appliance “Filers” – distributed file system appliances
- “Write-anywhere file layout”
- Serves up NFS, CIFS, http, ftp
- Random I/O optimized, write optimized
 - NVRAM for write caching
- Similar to Berkeley Fast File System, with extensive modifications


The WAFL File Layout




Snapshots in WAFL



(a) Before a snapshot.



(b) After a snapshot, before any blocks change.



(c) After block D has changed to D'.

Thanks



**Politecnico
di Torino**

Department of Control and
Computer Engineering



Questions?

sarah.azimi@polito.it